

Alexander Chechenev

AI Engineer (LLM, RAG, evaluation, agents)

Remote, English B2

alexchea319@gmail.com t.me/alexchea318 linkedin.com/in/alexander-chechenev

Summary

4+ years in software engineering, building LLM products since 2023. I lead development across all 7 services of an enterprise RAG platform, shipping cross-service features from technical plan to release. 3 on-prem deployments at 89–96% answer accuracy, measured by LLM-as-a-Judge.

Experience

Just AI (Jay Knowledge Hub)

Apr 2024 – present

Enterprise RAG platform in the Russian software registry, deployed in banking and manufacturing.

AI Engineer (RAG) Dec 2023 – present

- Built quality evaluation and a CI regression gate: test sets, LLM-as-a-Judge. Assistant coverage grew from 55 to 3,000 questions, accuracy holding at 89%. `Python`, `pytest`, `CI`
- Shipped document versioning across all 7 services: version became a chunk attribute, with no reindexing and no breaking API changes. `Kotlin`, `Python`, `TypeScript`
- Own on-prem delivery targeting 10,000 users at 1 RPS per user: LTS branches and rollout ordering across shared packages. `Docker`, `Kubernetes`, `Jenkins`
- Eliminated 3 sources of unnecessary spend: duplicate generation (semantic cache), duplicate crawler charges, billing during outages. Resolved 2 P1 incidents the same day. `Memcached`, `Grafana`, `Sentry`
- Own retrieval: hybrid search, query transformation, reranking, GraphRAG. Moved permission checks into search itself, closing 5 bypasses across 4 services. `Elasticsearch`, `Neo4j`, `RBAC`
- Codified the development cycle into runbooks executable by humans and AI agents: 2 engineers now deliver what previously required 3. Built and published the knowledge-base MCP server (5 tools). `Claude Code`, `MCP`, `GitLab CI`

Senior Frontend Developer Apr 2024 – Dec 2023

- Owned the frontend architecture: 4 repositories, 12 developers, zero audit vulnerabilities. Built the chat interface with streaming answers and citations. `React`, `Next.js`, `TypeScript`
- Led the migration of a 12-project monorepo (`React 16` → `19`, `webpack` → `Vite`, `Jest` → `Vitest`); 3 projects load 2.7× faster. `Vite`, `Vitest`, `Node.js`

NeoBIT, Development Team Lead and PM

Sep 2023 – Apr 2024

Generative AI for social media and messengers; Go services on Kubernetes.

- Led a team of 5 engineers across Go, frontend, DevOps and ML, taking 3 generative-AI products from requirements to production across 17 sprints, all on time.
- Combined Team Lead and PM responsibilities: 15 epics from business requirements, 8 demos to customer executives, 3 air-gapped deliveries for government contracts.

Earlier roles

NeoBIT, Frontend and Fullstack Developer. 2022–2023: `React`, a component library, the first Go services.

LG Electronics Russia R&D Lab, Data Science Intern. Summer 2021: optimization methods for gradient descent.

Stack

LLM and RAG: agents, tool calling, MCP, hybrid search, reranking, GraphRAG, `Elasticsearch`, `Neo4j`, `LlamaIndex`

Evaluation: LLM-as-a-Judge, test sets, prompt regression, quality metrics in CI, cost and latency

Engineering: `Python`, `FastAPI`, `Kotlin`, `Go`, `TypeScript`, `React`, `Next.js`, `PostgreSQL`, `MinIO`, `Memcached`

Delivery: `Docker`, `Kubernetes`, `Jenkins`, `GitLab CI`, on-prem, `Prometheus`, `Grafana`, `Sentry`

Education and talks

Specialist degree in Information and Analytical Security Systems, Peter the Great St. Petersburg Polytechnic University

2025: talk “Your own ChatGPT in 2021 for the price of a joke parser”, Elagin Pro festival